

KnowledgeHub — Project Documentation

Designing a Scalable Knowledge Base System System Design • Semester 4

This document is the source for the consolidated PDF [KnowledgeHub_Documentation.pdf](#) (build it with [python scripts/build_pdf.py](#)). A companion **Explanation Guide** ([KnowledgeHub_Explanation_Guide.pdf](#)) walks through every file and the system-design concepts in plain language.

Table of Contents

1. Problem Statement
 2. Proposed Solution
 3. Requirements Analysis (Q1)
 4. System Architecture (Q2)
 5. Document Management & Search Workflow (Q3)
 6. Database Design (Q4)
 7. Indexing & Search Algorithm (Q5)
 8. Scalability & Fault Tolerance (Q6)
 9. Module Description
 10. Technology Stack
 11. Implementation Details
 12. Screenshots
 13. Comparison with Real-World Systems
 14. Future Scope
-

1. Problem Statement

KnowledgeHub is an enterprise knowledge management platform — comparable to Confluence and Notion — that enables organizations to create, organize, search, and share technical documents, policies, FAQs, project documentation, and internal resources across teams. Employees use web and mobile clients to collaborate on documents, retrieve information quickly, and maintain organizational knowledge.

The platform must support **millions of documents and concurrent users** while delivering **fast search, collaboration features**, and **secure access control**. The core engineering challenges are:

- distributed document storage,
- maintaining version consistency,
- full-text indexing at scale,
- managing fine-grained user permissions, and
- enabling scalable collaboration workflows,

all while guaranteeing **high availability, fault tolerance, efficient synchronization, and reliable backups**.

2. Proposed Solution

We propose a **service-oriented, horizontally scalable architecture** with a clear separation between the **system of record** and **derived read models**:

- A **stateless FastAPI application tier** exposes authentication, document, search, and collaboration APIs and can be replicated behind a load balancer.
- **PostgreSQL** is the single source of truth for documents, immutable version history, permissions (ACLs), comments, and audit logs.
- **Elasticsearch** maintains a full-text inverted index for sub-second search; it is a *derived* store that can be rebuilt from PostgreSQL.
- **Redis** caches hot document reads and permission-scoped search results to absorb read traffic.
- An **indexing pipeline** keeps the search index synchronized whenever documents change.
- **Object storage (S3-compatible)** holds large attachments/blobs (design-level; the reference implementation stores text in Postgres).

A defining trait of the implementation is **graceful degradation**: each external dependency (Elasticsearch, Redis, Postgres) has a built-in fallback (a pure-Python inverted index, an in-memory cache, and SQLite respectively), so the system remains functional and demonstrable even when infrastructure is unavailable.

3. Requirements Analysis (Q1)

3.1 Functional Requirements

#	Requirement
F1	Users can register, authenticate, and receive a session token.
F2	Users can create, read, update, and delete documents, organised into spaces and tagged.
F3	Every document edit is preserved as an immutable version (revision history, diff/restore capable).
F4	Users can perform full-text keyword search and receive ranked, highlighted results.
F5	Owners can share documents and grant per-document access levels (owner/editor/viewer).
F6	Users can collaborate via comments and view a per-document activity/audit feed.
F7	Search results respect permissions — users only see documents they are allowed to access.

3.2 Non-Functional Requirements

#	Requirement	Target
N1	Search latency	p95 < 200 ms for typical queries
N2	Scalability	Millions of documents; tens of thousands of concurrent users
N3	Availability	99.9%+ via stateless tier + replicated data stores
N4	Consistency	Strong consistency for document content/versions; eventual for the search index
N5	Security	Encrypted transport, hashed passwords (bcrypt), JWT, RBAC + per-doc ACL
N6	Durability	No data loss on node failure; point-in-time backups
N7	Maintainability	Modular services, clear API contract, automated tests

3.3 Why search efficiency, collaboration, and scalability matter

- **Search efficiency** is the product. A knowledge base that can't surface the right document in milliseconds is functionally useless; employees abandon it and knowledge silos re-form. An **inverted index** turns an $O(N \cdot \text{document-size})$ scan into an $O(\text{query-terms})$ lookup, which is the only way to stay fast as the corpus grows to millions of documents.
- **Collaboration support** is what keeps knowledge *fresh*. Versioning, comments, and activity feeds let many people safely co-own a document, review changes, and trust that what they read is current — the difference between a living wiki and a graveyard of stale pages.
- **Scalability** is non-negotiable at enterprise size. Read and write volume, index size, and concurrent editing all grow with the organization. A **stateless application tier** (scale by adding replicas) plus **independently scalable data stores** (Postgres read replicas, Elasticsearch shards, Redis) lets each dimension scale on its own curve without rearchitecting.

4. System Architecture (Q2)

The platform is layered into **clients** → **gateway** → **stateless API services** → **data tier**. The editable Mermaid sources for every diagram below live in `docs/diagrams/` and `docs/architecture.md`.

Simple overview (the big picture in one line)





Figure 1 — At its simplest: a user works in the Web App, which calls the KnowledgeHub API over REST (secured with JWT). The API uses three data stores — PostgreSQL as the system of record, Elasticsearch for fast full-text search, and Redis to cache hot results.

Layered system architecture (presentation view)






Figure 2 — The same system in four clear layers, easy to narrate top-to-bottom: (1) users on the web/mobile app, (2) the gateway that secures and load-balances requests, (3) the stateless API and its four responsibilities, and (4) the three data stores, each doing one job — PostgreSQL (source of truth), Elasticsearch (fast search), and Redis (cache). A fully detailed component diagram (with the indexing pipeline, object storage and backups) is available at docs/diagrams/01_architecture.png.

Write path (create / update a document)




Figure 3 — A document edit is committed to PostgreSQL (row + immutable version + audit log), then projected into the Elasticsearch index, and the affected caches are invalidated.

Search path (query)




Figure 4 — Search first consults the per-user Redis cache; on a miss it resolves the caller's visible document IDs from the ACL, queries Elasticsearch filtered to those IDs, then caches the result.

Component responsibilities

- **API Gateway / Load Balancer** — TLS termination, rate limiting, routing, and fan-out across stateless API replicas.
- **Auth Service** — registration, login (JWT), and authorization (global RBAC + per-document ACL resolution).
- **Document Service** — CRUD, versioning, and orchestration of index sync + cache invalidation.
- **Search Service** — resolves the caller's visible document set, queries the index, caches results.
- **Collaboration Service** — comments and the append-only activity feed.
- **Indexing Pipeline** — projects document changes into the search index (synchronous in the reference build; queue-based in production).
- **PostgreSQL / Elasticsearch / Redis / Object Storage** — the data tier described above.

5. Document Management & Search Workflow (Q3)

5.1 Creation & versioning

1. An authenticated user **POST**s a document. The Document Service writes the **documents** row and a **document_versions** row (version 1) in one transaction.
2. On every update, the service compares fields, increments **current_version**, and appends a **new immutable version row** — history is never overwritten. This makes diffs, audits, and rollbacks trivial.

5.2 Indexing

3. After a successful commit, the **indexing pipeline** upserts **{title, content, tags, owner, space}** into Elasticsearch. The relational store stays authoritative; the index is a rebuildable projection.

5.3 Search

4. A query first checks **Redis** using a **per-user cache key** (so permission-filtered results never leak between users).
5. On a miss, the Search Service resolves the caller's **visible document IDs** from the ACL, queries Elasticsearch **filtered to those IDs**, and ranks results (TF-IDF/BM25) with highlights.
6. The result is cached with a TTL and returned. Document edits **invalidate** the relevant cache entries.

5.4 Synchronization & sharing

7. Sharing writes a **permissions** (ACL) row; authorization is enforced **at the data layer**, not just the UI, so search and direct reads both honour it.
8. Every action (**created/updated/shared/commented/deleted**) appends to **collaboration_logs**, powering activity feeds and providing an audit trail for synchronization across distributed teams.

Security across distributed teams is enforced by: TLS in transit, bcrypt-hashed credentials, short-lived JWTs, global roles plus per-document ACLs, and permission-filtered search.

6. Database Design (Q4)

We use a **relational schema (PostgreSQL)** as the system of record because documents, versions, permissions, and logs are highly relational and benefit from ACID guarantees and foreign-key integrity. Full-text data is offloaded to **Elasticsearch**

(a NoSQL document/search store) where inverted-index search excels — a deliberate **polyglot-persistence** split: SQL for truth + relationships, NoSQL for search.

6.1 Entity-Relationship overview




Figure 5 — Relational schema (PostgreSQL). **document_versions** is append-only for full revision history; **permissions** is the per-document ACL; **collaboration_logs** is the append-only audit trail.

6.2 Tables & rationale

Table	Purpose	Key design notes
<code>users</code>	Accounts & global role	Unique index on <code>email</code> ; bcrypt hash stored, never the password.
<code>spaces</code>	Team workspaces	Groups related documents for navigation and bulk permissions.
<code>documents</code>	Current document state	Indexes on <code>title</code> , <code>owner_id</code> , <code>space_id</code> ; <code>current_version</code> pointer; denormalised <code>tags</code> for fast filtering.
<code>document_versions</code>	Immutable revision history	<code>UNIQUE(document_id, version)</code> ; append-only — enables diff/restore/audit.
<code>permissions</code>	Per-document ACL	<code>UNIQUE(document_id, user_id)</code> ; level ∈ {owner, editor, viewer}.
<code>comments</code>	Collaboration threads	Indexed by <code>document_id</code> .
<code>collaboration_logs</code>	Append-only audit/activity	Indexed by <code>document_id</code> and <code>created_at</code> for fast feeds.

6.3 Elasticsearch document mapping (search store)

```
{
  "title": { "type": "text", "analyzer": "english", "fields": { "raw": { "type": "keyword" } } },
  "content": { "type": "text", "analyzer": "english" },
  "tags": { "type": "keyword" },
  "space_id": { "type": "keyword" },
  "owner_id": { "type": "keyword" },
  "updated_at": { "type": "date" }
}
```

Sharding/replication of this index is what lets search scale horizontally and stay highly available.

7. Indexing & Search Algorithm (Q5)

The reference algorithm lives in `app/services/inverted_index.py` (standard library only) and is demonstrated in `notebooks/demo.ipynb` .

7.1 Pipeline

1. **Analysis** — lower-case, tokenize on word boundaries, remove stop-words. The *same* analyzer is applied to documents and queries so terms match consistently.
2. **Indexing** — build a postings map `term -> {doc_id: term_frequency}` and track each document's length.
3. **Weighting** — at query time compute `idf(t) = ln(1 + N / df(t))` and TF-IDF weights per (doc, term).
4. **Ranking** — score documents by **cosine similarity** between the query and document TF-IDF vectors, normalising by vector magnitude so long documents aren't unfairly favoured.

7.2 Core logic (excerpt)

```
def search(self, query, limit=10):
    q_tokens = tokenize(query)
    q_tf = Counter(q_tokens)
    q_vec = {t: tf * self._idf(t) for t, tf in q_tf.items()}
    q_norm = sqrt(sum(w*w for w in q_vec.values())) or 1.0

    scores = defaultdict(float)
    for term, q_weight in q_vec.items():
        idf = self._idf(term)
        for doc_id, tf in self.postings.get(term, {}).items():
            d_weight = (tf / self.doc_length[doc_id]) * idf  # TF normalised by length
            scores[doc_id] += q_weight * d_weight

    ranked = [(d, dot / (q_norm * self._doc_norm(d)), self._snippet(d, q_tokens))
               for d, dot in scores.items()]
    ranked.sort(key=lambda x: x[1], reverse=True)
    return ranked[:limit]
```

7.3 Why this design

- **Inverted index** → search touches only documents that contain a query term, not the whole corpus → scales to millions of documents.
- **IDF** down-weights ubiquitous words and rewards rare, discriminating terms.
- **Cosine similarity** gives a length-fair relevance score in [0, 1].
- These are exactly the principles behind Lucene/Elasticsearch (which use the closely related **BM25**); the reference code makes the mechanics explicit and serves as the offline fallback engine.

7.4 Worked example (from the notebook)

```
Query: 'production deployment'
0.4749 Kubernetes Production Deployment Guide      Kubernetes deployment guide for production clu
0.4749 Deploy Checklist                            Production deployment checklist and rollback p
```

8. Scalability & Fault Tolerance (Q6)

8.1 Scaling to millions of documents & users

Dimension	Strategy
Application tier	Stateless API → add replicas behind the load balancer; no sticky sessions (JWT).
Search	Elasticsearch sharding distributes the index across nodes; replicas add read throughput and HA.
Database reads	PostgreSQL read replicas serve heavy read traffic; primary handles writes.
Database writes / data size	Partition/shard large tables (e.g. by space or tenant); archive cold versions.
Caching	Redis absorbs repeated reads and search queries, cutting load on Postgres/ES.
Async indexing	A message queue (e.g. Kafka/RabbitMQ) decouples writes from indexing so write latency is unaffected by index load.
Object storage	Large attachments stream from S3/CDN, keeping the database lean.

8.2 Fault tolerance & recovery

- **Stateless services** — a failed API pod is simply replaced; the load balancer routes around it.
 - **Replication** — Postgres streaming replication with automatic failover; Elasticsearch replica shards; Redis replicas/Sentinel.
 - **Source of truth + rebuildable index** — if Elasticsearch data is lost, the index is **rebuilt from PostgreSQL**; no permanent data loss.
 - **Graceful degradation (implemented here)** — if Elasticsearch is down, search falls back to the local inverted index; if Redis is down, the cache falls back to in-memory; if Postgres is unavailable, the dev/demo profile uses SQLite.
`GET /health` reports the active backends.
 - **Backups** — automated Postgres snapshots / PITR (point-in-time recovery) and periodic Elasticsearch snapshots to object storage.
 - **Idempotent indexing** — upserts keyed by document ID make retries safe after partial failures.
-

9. Module Description

Module	File(s)	Responsibility
Configuration	<code>app/config.py</code>	Environment-driven settings with safe local fallbacks.
Persistence	<code>app/database.py</code> , <code>app/models.py</code>	SQLAlchemy engine/session and the relational schema.
Schemas	<code>app/schemas.py</code>	Pydantic request/response contracts and validation.
Authentication & Authorization	<code>app/auth.py</code>	bcrypt hashing, JWT issue/verify, RBAC + per-doc ACL checks.
Cache	<code>app/cache.py</code>	Redis client with transparent in-memory fallback; namespace invalidation.
Search	<code>app/search.py</code>	Elasticsearch lifecycle + query; inverted-index fallback.
Inverted Index	<code>app/services/inverted_index.py</code>	TF-IDF + cosine keyword search (Q5).
Document Service	<code>app/services/document_service.py</code>	Versioning, index sync, cache invalidation, audit logging.
API Routers	<code>app/routers/*.py</code>	HTTP endpoints for auth, documents, search, collaboration.
App entrypoint	<code>app/main.py</code>	Wires routers, mounts the Web UI, creates tables, exposes <code>/health</code> .
Web UI (frontend)	<code>app/static/index.html</code>	Single-page app (login, create, search, view, comment) served at <code>/app</code> .
Demo & tests	<code>notebooks/demo.ipynb</code> , <code>smoke_test.py</code>	Executable demonstration and end-to-end verification.

10. Technology Stack

Layer	Technology	Why
Frontend	Vanilla HTML/CSS/JS (single-page)	Zero-build UI served by the API; talks to REST endpoints with JWT.
API framework	FastAPI (Python)	Async, type-safe, auto-generated OpenAPI/Swagger docs.
Server	Uvicorn	High-performance ASGI server.
ORM	SQLAlchemy 2.0	Mature, DB-agnostic (Postgres in prod, SQLite for demo).
Relational DB	PostgreSQL	ACID source of truth, strong relational integrity.
Search engine	Elasticsearch	Distributed full-text search, sharding, BM25 ranking.
Cache	Redis	Low-latency caching of reads and search results.
Auth	JWT (python-jose) + bcrypt	Stateless auth + secure password hashing.
Packaging	Docker / Docker Compose	Reproducible multi-service stack.
Validation	Pydantic v2	Request/response validation and serialization.

11. Implementation Details

- **Versioning** is implemented by appending to `document_versions` on every content change while bumping `documents.current_version`; the unique constraint `(document_id, version)` guarantees a clean history.
- **Authorization** resolves an *effective level* per (user, document): admins and owners get `owner`; otherwise the ACL entry applies. `require_access` enforces a minimum level on every protected endpoint and is reused by search filtering.
- **Permission-filtered search** computes the caller's visible document IDs and passes them as an Elasticsearch `ids` filter (or filters the fallback results), so authorization is enforced in the data path — not just hidden in the UI.
- **Cache correctness** uses **per-user, query-hashed keys** for search results and `doc:{id}` keys for reads; writes invalidate both, preventing stale or cross-user leakage.
- **Graceful degradation** is handled at construction time: the cache and search modules probe their backends and silently switch to fallbacks, logging which backend is active (surfaced via `/health`).
- **Verification:** `smoke_test.py` exercises register → login → create → version → permission-deny → share → permission-allow → search (filtered) → cache hit → comment → activity feed, and asserts each invariant. The notebook reproduces this interactively with embedded outputs.

12. Screenshots

All screenshots were captured from the **live dockerized stack** (PostgreSQL + Elasticsearch + Redis).

Web application (the user-facing UI)

The single-page Web App is served by the API at `http://localhost:8001/app`.

Sign-in screen

KnowledgeHub

Enterprise knowledge base — documents, search & collaboration

Sign in

Register

Email

alice@corp.com

Password

Sign in

Tip: a demo user `alice@corp.com` / `secret1` may already exist.

Figure 6 — Login / registration (OAuth2 password flow → JWT).

Dashboard — create & search

KnowledgeHub

search: elasticsearch

Alice Nguyen (editor)

Sign out

Create document

Title

e.g. Kubernetes Deployment Guide

Content

Write the document body...

Tags (comma separated)

devops, k8s

Create & index

Search

production deployment

Go

3 result(s) · 68 ms · fresh

score 5.837

Kubernetes Production Deployment Guide

This guide explains how to **deploy** services to **production** Kubernetes clusters.

score 0.489

Incident Response Runbook

Steps to triage, contain and remediate a **production** incident: declare severity, assemble the response

score 0.456

Disaster Recovery Policy

and RTO targets, cross-region replication, and the step-by-step database restore procedure for **production**

Select a search result to open a document.

Figure 7 — Creating and searching documents. Results show the live Elasticsearch BM25 score, matched-term highlights, query latency, and whether the result came from the Redis cache. The badge confirms `search: elasticsearch` .

Document view — content, version history & comments




Figure 8 — Opening a document shows its content, tags, full revision history (v2, v1), and the collaboration comment thread.

API & infrastructure

Interactive API documentation (Swagger UI) — <http://localhost:8001/docs>

KnowledgeHub

1.0.0 OAS 3.1

[/openapi.json](#)

Enterprise knowledge management platform — documents, search, collaboration.

Authorize

auth

POST /auth/register Register

POST /auth/login Login

GET /auth/me Me

documents

POST /documents Create

GET /documents/{doc_id} Read

PUT /documents/{doc_id} Update

DELETE /documents/{doc_id} Delete

GET /documents/{doc_id}/versions Versions

POST /documents/{doc_id}/share Share

collaboration

POST /documents/{doc_id}/comments Add Comment

GET /documents/{doc_id}/comments List Comments

GET /documents/{doc_id}/activity Activity Feed

search

GET /search Search

system

GET / Root

GET /health Health

Schemas

Body_login_auth_login_post > Expand all object

CommentCreate > Expand all object

Figure 9 — Auto-generated OpenAPI documentation listing every endpoint grouped by module (auth, documents, collaboration, search, system).

Health check — active backends — <http://localhost:8001/health>

```
Pretty print ☐
{"status":"ok","search_backend":"elasticsearch","cache_backend":"Redis"}
```

Figure 10 — The running platform reports `search_backend: elasticsearch` and `cache_backend: Redis` , confirming the full distributed stack is active.

End-to-end run on the real stack

```

KnowledgeHub — end-to-end demo (Postgres + Elasticsearch + Redis)

$ curl http://127.0.0.1:8001/health
{'status': 'ok', 'search_backend': 'elasticsearch', 'cache_backend': 'Redis'}

# Register an author (editor) and a teammate (viewer)
alice + bob registered; JWTs issued

# Alice creates documents -> persisted in Postgres, indexed in Elasticsearch
+ 7cf39b7a v1 Kubernetes Production Deployment Guide
+ 03b838e1 v1 Disaster Recovery Policy
+ 56ef0285 v1 Onboarding FAQ

# Versioning: editing creates a new immutable revision
revision history (latest first): [2, 1]

# Authorization: Bob cannot read the doc until it is shared
GET /documents/{k8s} as bob -> 403 (403 Forbidden)
POST /documents/{k8s}/share (bob: viewer)
GET /documents/{k8s} as bob -> 200 (200 OK)

# Full-text search via Elasticsearch (BM25 ranking + highlights)
q='production deployment' (2 hits, 1138ms, cached=False)
  3.800 Kubernetes Production Deployment Guide
    <em>Deploy</em> to <em>production</em> Kubernetes; adds blue-green and canary strategies.
  0.348 Disaster Recovery Policy
    Backup schedules, RPO and RTO targets, and the <em>production</em> database restore procedure.

# Permission-filtered search: Bob only sees documents shared with him
bob q='production' -> ['Kubernetes Production Deployment Guide']

# Caching: identical repeated query served from Redis
repeat query -> cached=True (0ms)

# Collaboration: comment + activity feed
[ commented] Should we add a Helm section?
[   shared] f16f1e04-40ec-4ac2-a3e9-0b4573449559:viewer
[  updated] v2
[   created] Created 'Kubernetes Production Deployment Guide'

All features verified on Postgres + Elasticsearch + Redis.

```

Figure 11 — A complete run (`scripts/demo_transcript.py`): registration + JWT, document creation with Elasticsearch indexing, versioning, `403` → `200` permission enforcement, BM25 search with highlights, permission-filtered results, Redis cache hit, and the collaboration activity feed.

13. Comparison with Real-World Systems

Aspect	KnowledgeHub	Confluence	Notion	Elasticsearch-backed wikis
Storage of truth	PostgreSQL	Relational DB (e.g. Postgres)	Block-based store	Varies
Search	Elasticsearch (BM25) + local fallback	Lucene/Elasticsearch	Internal search	Elasticsearch
Versioning	Immutable version rows	Page history	Page history / version log	Custom
Permissions	RBAC + per-doc ACL	Space + page restrictions	Workspace + page sharing	Custom
Architecture	Stateless services + polyglot persistence	Service-oriented	Cloud, block model	Index-centric

KnowledgeHub mirrors the **proven enterprise pattern**: a relational system of record plus a dedicated, independently scalable search index — the same separation Confluence (Lucene/ES) and large wiki deployments rely on. The contribution here is a compact, transparent reference implementation with explicit fallbacks and a from-scratch ranking algorithm for learning purposes.

14. Future Scope

- **Real-time co-editing** via WebSockets + CRDTs/OT (Notion/Google-Docs-style concurrent editing).
- **Asynchronous indexing** through Kafka/RabbitMQ for write/index decoupling at scale.
- **Semantic / vector search** using embeddings (e.g. dense-vector kNN in Elasticsearch) for natural-language retrieval and RAG-style Q&A over the knowledge base.
- **Rich content**: attachments to object storage, images, diagrams, and Markdown/WYSIWYG editing.
- **Diff & restore UI** for version history; merge-conflict handling.
- **Full-text highlighting & faceted search** by space, tag, author, and date.
- **Observability**: metrics, tracing, and per-tenant rate limiting.
- **Multi-tenancy & SSO** (SAML/OIDC) for enterprise deployments.

Appendix — Build the submission PDF

```
pandoc DOCUMENTATION.md docs/architecture.md \  
-o KnowledgeHub_Documentation.pdf \  
--toc --highlight-style=tango
```

(Or open `DOCUMENTATION.md` in VS Code with a Markdown-PDF extension / Typora and export.)